
2023 Spring RL Final Presentation

Batch Constraint Q-Learning

Members:

111550006 林庭寫

111550093 朱驛庭

111550106 何義翔

111550160 黃佑得





Table of contents

01

Introduction

The idea of BCQ

02

Reproduction

Reproduction of the
result on D4RL

03


Ablation Study

finding better nn settings
on gym task

04

Adroit Task

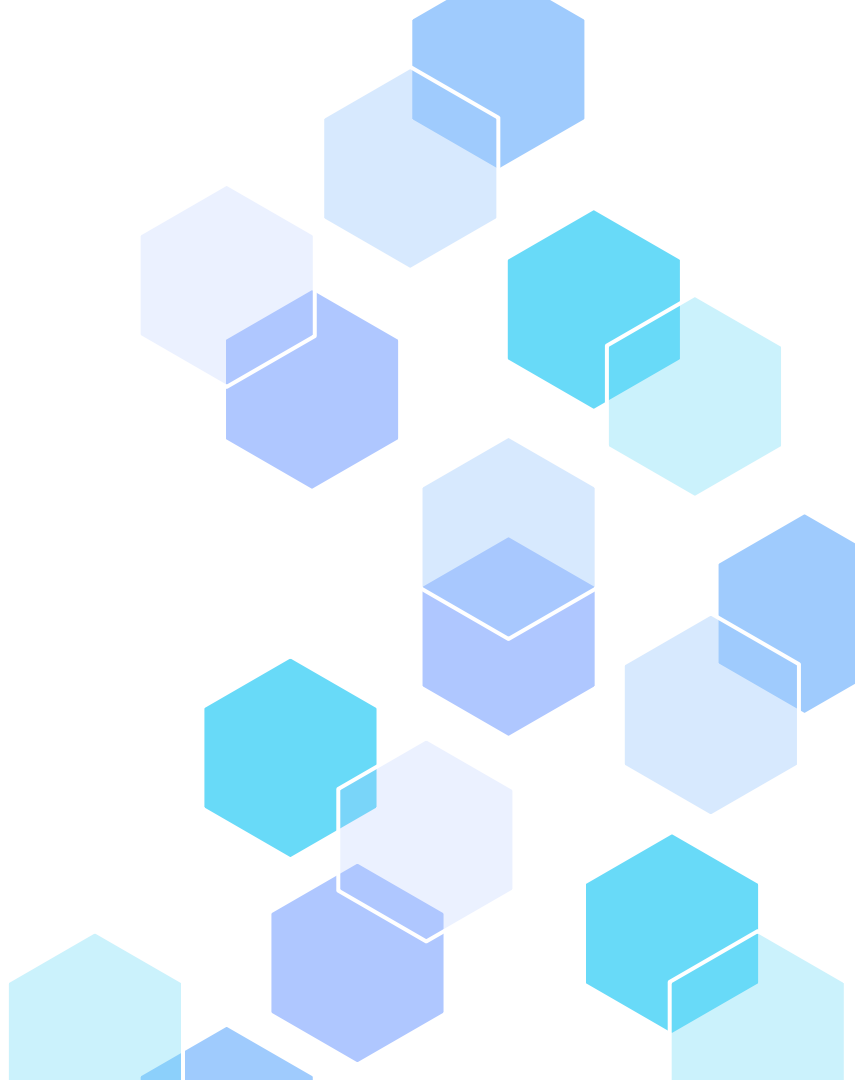
Try to improve returns in
Adroit tasks



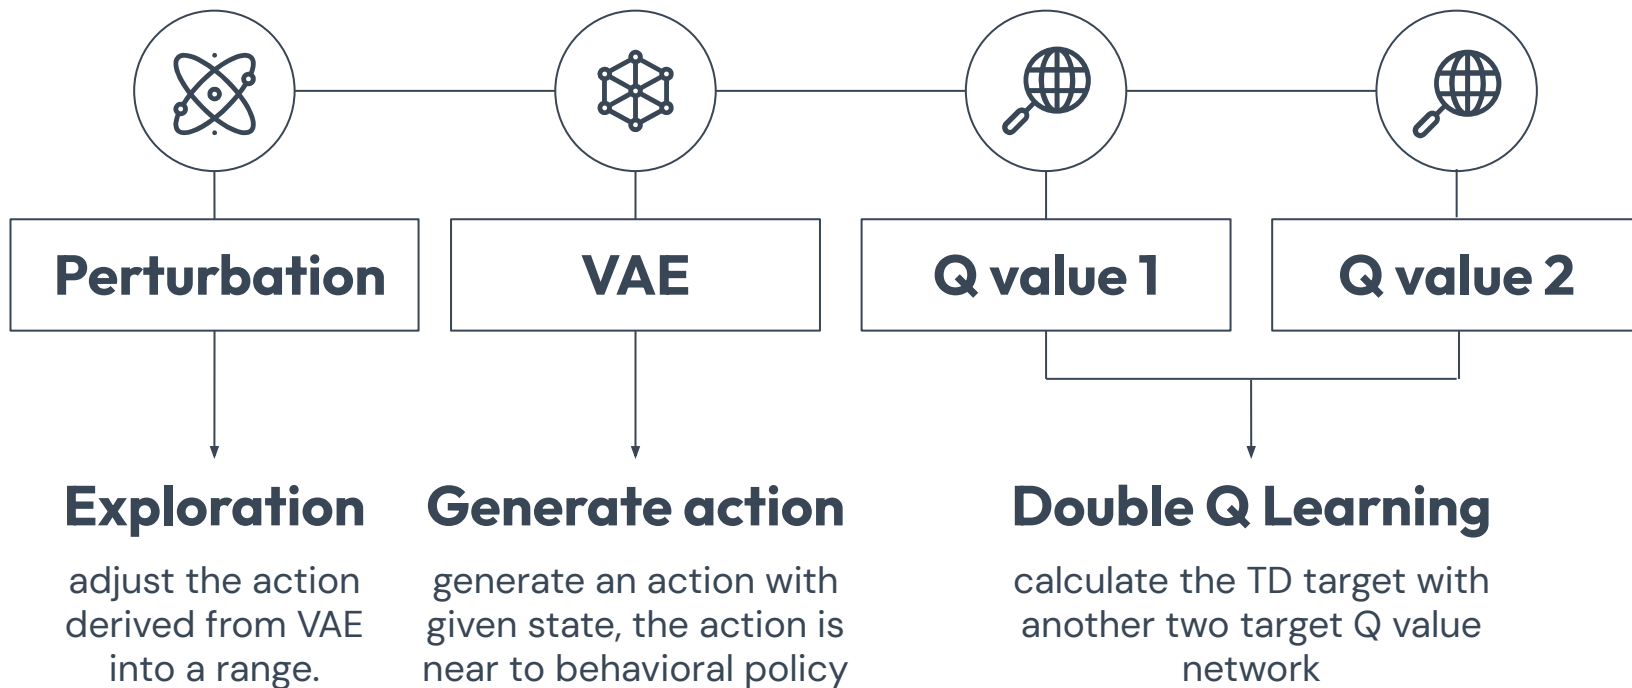
01

Introduction

Brief intro. to today's main method, BCQ



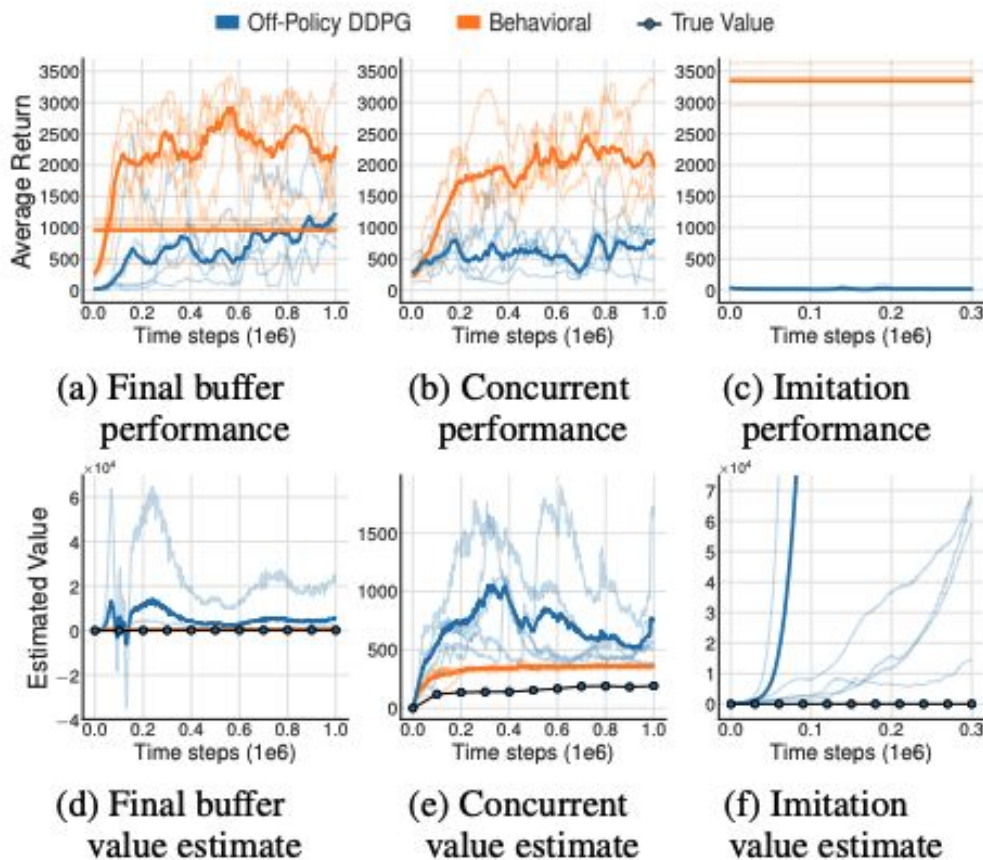
Four Networks in BCQ



Why BCQ?

Try to solve **Extrapolation error**, caused by:


1. Absent data
2. Model Bias
3. Training mismatch



02

Reproduction

Reproduce the result from original paper



Maze2D & Hopper Reproduce

		seed=123	seed=456	seed=789	avg	paper
Maze2D	ant-maze-medium-diverse-v0	0.00	0.00	0.00	0.00	0.00
	ant-maze-medium-play-v0	0.00	0.30	0.00	0.10	x
	maze2d-large-v1	109.10	119.10	90.20	106.13	23.20
	maze2d-medium_v1	85.60	71.10	72.60	76.43	35.00
	maze2d-umaze-v1	91.50	76.90	81.30	83.23	41.50

hopper	hopper-random-v0	330.58	324.36	322.87	325.94	323.90
	hopper-medium-v0	929.84	2664.16	988.39	1527.46	1752.40
	hopper-expert-v0	3679.69	2793.27	3664.08	3379.02	x
	hopper-medium-replay-v0	1185.01	726.79	1165.64	1025.81	1057.80
	hopper-medium-expert-v0	3667.69	3518.68	3668.15	3618.17	3588.50

Halfcheetah & Walker2D Reproduce

		seed=123	seed=456	seed=789	avg	paper
halfcheetah	halfcheetah-random-v0	-1.12	-1.21	-1.46	-1.26	-1.30
	halfcheetah-medium-v0	4688.00	4462.00	4721.00	4623.67	4767.90
	halfcheetah-expert-v0	11211.00	10806.00	12307.00	11441.33	x
	halfch.-medium-replay-v0	4371.00	4315.00	4125.00	4270.33	4463.90
	halfch.-medium-expert-v0	12073.00	10727.00	12615.00	11805.00	7750.80

walker2d	walker2d-random-v0	220.94	221.52	222.11	221.52	228.00
	walker2d-medium-v0	2815.42	2847.64	2164.88	2609.31	2441.00
	walker2d-expert-v0	4549.01	3366.86	4521.77	4145.88	x
	walker2d-medium-replay-v0	716.67	779.04	822.28	772.66	688.70
	walker2d-medium-expert-v0	3413.47	3823.71	3335.65	3524.28	2640.30

03

Ablation Study

Find better hyperparameters for gym task.





hopper-random-v0

The reward is relatively stable

Iterations = 1e5


Quick convergence of hopper-random-v0

Seed = 123, 456, 789

Get the average reward of three seeds

Multiple hidden layers

Rewards



Deeper VAE & Q Network

Multiple VAE hidden layers

The performance maintains .



Follow the link in the graph to modify its data and then paste the new one here. [For more info, click.](#)

CQL + BCQ

$$\alpha \mathbb{E}_{s \sim \mathcal{D}} [\log \sum_a \exp(Q(s, a)) - \mathbb{E}_{a \sim \hat{\pi}_\beta(a|s)} [Q(s, a)]] + \frac{1}{2} \mathbb{E}_{s, a, s' \sim \mathcal{D}} \left[\left(Q - \hat{\mathcal{B}}^{\pi_k} \hat{Q}^k \right)^2 \right]$$

$$\log \sum_a \exp(Q(s, a))$$

Logsumexp

Try to approximate
the original bound

$$\mathbb{E}_{a \sim \hat{\pi}_\beta(a|s)} [Q(s, a)]$$

Current Q

Directly use current
Q-value

$$\mathbb{E}_{s, a, s' \sim \mathcal{D}} \left[\left(Q - \hat{\mathcal{B}}^{\pi_k} \hat{Q}^k \right)^2 \right]$$

Critic Loss

The Bellman error

CQL + BCQ



```
current_Q1, current_Q2 = self.critic(state, action)
vae_Q1, vae_Q2         = self.critic(state, self.vae.decode(next_state))
pred_Q1, pred_Q2       = self.critic(state, self.actor.forward(state, action))
```

[For reference, click.](#)

CQL + BCQ




```
alpha = torch.tensor(0.4, dtype=torch.float32)
q1_cat = torch.cat([current_Q1, vae_Q1, pred_Q1], 1)
q2_cat = torch.cat([current_Q2, vae_Q2, pred_Q2], 1)

logsumexp_Q1 = torch.logsumexp(q1_cat, dim=1).mean()
logsumexp_Q2 = torch.logsumexp(q2_cat, dim=1).mean()

penalty1 = alpha*(logsumexp_Q1 - current_Q1)
penalty2 = alpha*(logsumexp_Q2 - current_Q2)

bellman_er1 = F.mse_loss(current_Q1, target_Q)
bellman_er2 = F.mse_loss(current_Q2, target_Q)
critic_loss = (0.5 * bellman_er1 + penalty1) + (0.5 * bellman_er2 +
penalty2)
```

Scaling Hidden Size




Delete Target Perturbation Network

Use the same
network to
evaluate Q
target

The performance
maintains.

Delete actor_target



Follow the link in the graph to modify its data and then paste the new one here. [For more info, click.](#)

Delete Target Perturbation Network



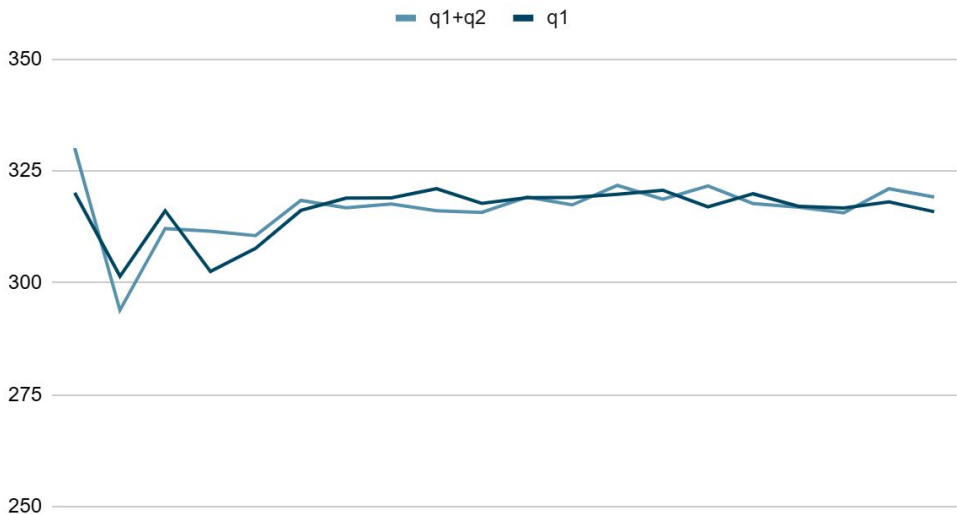
```
def train(self, replay_buffer, iterations, batch_size=100):  
    # ...  
    # Critic Training  
    with torch.no_grad():  
        # Duplicate next state 10 times  
        next_state = torch.repeat_interleave(next_state, 10, 0)  
        # Compute value of perturbed actions sampled from the VAE  
        target_Q1, target_Q2 = self.critic_target(next_state, self.actor(next_state, self.vae.decode(next_state)))  
    # ...
```

Use ($q1 + q2$) to Select Action

**The original
implement
use only $q1$**

The performance
maintains.

Select action



Follow the link in the graph to modify its data and then paste the new one here. [For more info, click.](#)

Use (q1 + q2) to Select Action

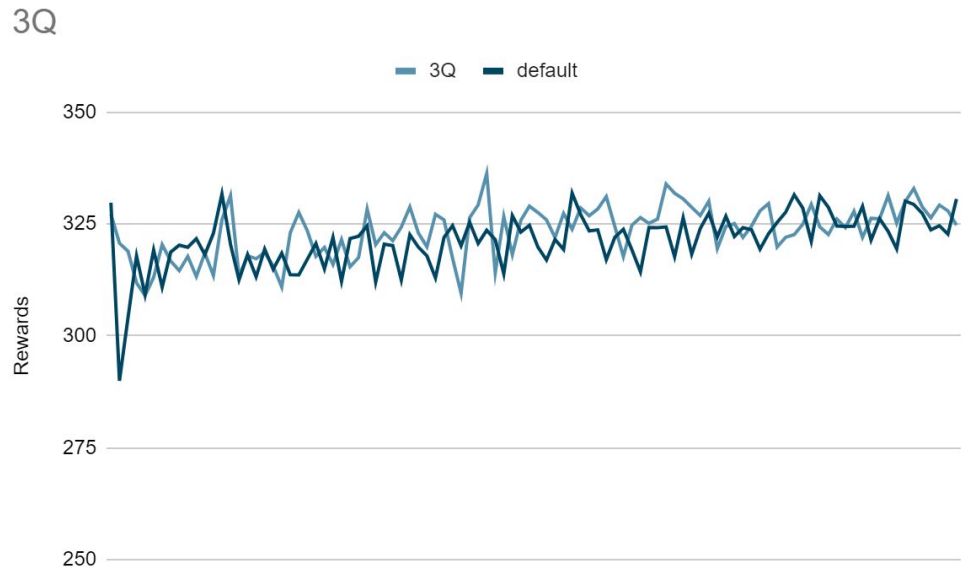
```
def select_action(self, state):  
    with torch.no_grad():  
        state = torch.FloatTensor(state.reshape(1, -1)).repeat(100, 1).to(self.device)  
        action = self.actor(state, self.vae.decode(state))  
        q1 = self.critic.q1(state, action)  
        ind = q1.argmax(0)  
    return action[ind].cpu().data.numpy().flatten()
```

Triple Q Networks

Use

3 Q-Networks

The performance maintains .



Follow the link in the graph to modify its data and then paste the new one here. [For more info, click.](#)

Triple Q Networks

```
# Critic Training
with torch.no_grad():
    # Duplicate next state 10 times
    next_state = torch.repeat_interleave(next_state, 10, 0)

    # Compute value of perturbed actions sampled from the VAE
    target_Q1, target_Q2, target_Q3 = self.critic_target(next_state, self.actor_target(next_state, self.vae.decode(next_state)))

    # Soft Clipped Double Q-learning
    #mini = torch.min(target_Q1, target_Q2)
    #maxi = torch.max(target_Q1, target_Q2)
    mini = torch.min(torch.min(target_Q1, target_Q2), target_Q3)
    maxi = torch.max(torch.max(target_Q1, target_Q2), target_Q3)
    target_Q = self.lmbda * mini + (1. - self.lmbda) * maxi
    # Take max over each action sampled from the VAE
    target_Q = target_Q.reshape(batch_size, -1).max(1)[0].reshape(-1, 1)

    target_Q = reward + not_done * self.discount * target_Q


current_Q1, current_Q2, current_Q3 = self.critic(state, action)
critic_loss = F.mse_loss(current_Q1, target_Q) + F.mse_loss(current_Q2, target_Q) + F.mse_loss(current_Q3, target_Q)


self.critic_optimizer.zero_grad()
critic_loss.backward()
self.critic_optimizer.step()
```

04

Adroit Task

solve a more complex task with BCQ





door

door-[human/cloned/expert]

Iterations = 5e5

Seed = 123, 456, 789

Get the average reward of three seeds

Aroit Task Results

	3 VAE + 3 hidden	wide hidden	Triple Q	delete actor target	select action	CQL+BCQ	Origin
door -human	-57.12	-57.03	-59.36	-59.95	-59.81	-53.46	-56.6
door -cloned	-59.21	-56.26	-55.27	-57.52	-57.45		-56.3
door -expert			2948.32	2992.68	2972.49	2943.71	2850.7



Thanks!